

**Course Notes for**  
**CS 1501**  
**Algorithm Implementation**

**By**  
**John C. Ramirez**  
**Department of Computer Science**  
**University of Pittsburgh**



- These notes are intended for use by students in CS1501 at the University of Pittsburgh and no one else
- These notes are provided free of charge and may not be sold in any shape or form
- These notes are NOT a substitute for material covered during course lectures. If you miss a lecture, you should definitely obtain both these notes and notes written by a student who attended the lecture.
- Material from these notes is obtained from various sources, including, but not limited to, the following:
  - Algorithms in C++ by Robert Sedgewick
  - Algorithms, 4<sup>th</sup> Edition by Robert Sedgewick and Kevin Wayne
  - Introduction to Algorithms, by Cormen, Leiserson and Rivest
  - Various Java and C++ textbooks
  - Various online resources (see notes for specifics)




- Consider BST search for key  $K$ :
  - ▶ For each node  $T$  in the tree we have 4 possible results:
    - 1)  $T$  is empty (or sentinel node) indicating item not found.
    - 2)  $K$  matches  $T.key$  and item is found.
    - 3)  $K < T.key$  and we go to left child.
    - 4)  $K > T.key$  and we go to right child.
  - ▶ Consider now the same basic technique, but proceeding left or right based on the **current bit** within the key:
    - We still have a tree and 1) and 2) are the same
    - But 3) and 4) will be different.



- Call this tree a **Digital Search Tree (DST)**.
- DST search for key K:
  - ▶ For each node T in the tree we have 4 possible results:
    - 1) T is empty (or a sentinel node) indicating item not found.
    - 2) K matches T.key and item is found.
    - 3) **Current bit of K is a 0** and we go to left child.
    - 4) **Current bit of K is a 1** and we go to right child.
      - ▶ **Each time we move down in the tree we go to the next bit in the key.**
  - ▶ Look at example on next page. →



# Digital Search Trees



**010010 = 18**

**001001 = 9**

**001000 = 8**

**101000 = 40**

**101001 = 41**

**010000 = 16**

**010100 = 20**

**011100 = 28**

**010011 = 19**

Note:

Go left on 0 bit

Go right on 1 bit



- Run-times?

- ▶ Given  $N$  random keys, the height of a DST should **average  $O(\log_2 N)$** .
  - Think of it this way – if the keys are random, at each branch it should be equally likely that a key will have a 0 bit or a 1 bit.
    - Thus the tree should be well-balanced.
- ▶ In the **worst case**, we are bound by the **number of bits in the key** (say it is  $b$  number of bits).
  - So in a sense we can say that this tree has a constant run-time, if the number of bits in the key is a constant.
    - This is an improvement over the BST.



- But DSTs have **drawbacks**:
  - 1) **Data is not sorted.**
    - If we want sorted data, we would need to extract all of the data from the tree and sort it.
  - 2) **Bitwise operations are not always easy.**
    - Some languages do not provide for them at all, and for others it is costly.
  - 3) **What about keys of different lengths?**
    - What if a key is a prefix of another key that is already present?
      - Where would we put it?



- 4) In addition to our bitwise comparisons, we are **still doing comparisons of the entire key**.
- Note that we **use the bits to branch**, but we **use the whole key for the equality test**.
  - May do  $b$  key comparisons in the worst case!
  - If a key is long and comparisons are costly, this can be inefficient.
  - BST does key comparisons as well, but does not have the additional bit comparison.
- All of these issues make DSTs less than desirable as an implementation of a symbol table.
- But maybe we can improve them...





- Let's first address the last problem (4):
  - How to reduce / eliminate the number of comparisons (of the entire key)?
    - This will also resolve (1) and (3).
    - We will resolve (2) soon.
  - We'll modify our tree slightly.
    - All values will be in exterior nodes at leaves in the tree.
    - Interior nodes will not contain keys, but will just direct us down the tree toward the leaves, based on the current bit value.
    - In fact, we do not need to store the keys themselves at all.




## Radix Search Tries

- Each path from the root to a leaf represents a unique key.
  - Thus the keys are “stored” within the trie, but implicitly via the paths to the leaves rather than explicitly.
- ▶ This gives us a **Radix Search Trie**:
  - Trie is from re**TRIE**val (see text).
  - See example with same data as DST on next slide. →
  - Note that the **tree is taller because all bits in the keys are used**.
  - Note also that the **keys are being stored in the leaves**.
    - This is assuming the symbol table is storing (key, key) pairs.
    - Normally, the symbol table will have (key, value) pairs where the value differs from the key.



# Radix Search Tries



- Benefit of simple Radix Search Tries
  - ▶ No comparisons of the **entire key**
    - Thus solves problem (4) of DSTs
  - ▶ Do **not** require **keys** to be **stored**
    - Paths from root implicitly define the keys
  - ▶ Can handle keys with common prefixes
    - Simply allow interior nodes to have values as well
      - Put value in next node after last bit of key to indicate the termination of a key
      - Thus solves problem (3) of DSTs



- ▶ An “inorder” traversal of the tree will give the data in “sorted” order (problem (1)).
  - Look again at the tree in slide 11 to see this.
- Drawbacks:
  - ▶ The tree will have **more overall nodes** than a DST.
    - Bits to each key are fully elaborated as branches in the tree.
    - Thus, we need to traverse  $b+1$  levels of the tree to find a key of  $b$  bits, regardless of the number of keys in the tree.
  - ▶ Many nodes will have an **unused value** reference.
    - Every node has this reference, but most will not store a value.
  - ▶ We are still doing **bit operations** (problem 2).



- Asymptotic run-time is similar to DST
  - Since all bits are elaborated, every **successful** search requires  $b$  bit comparisons where  $b$  is the number of bits in the key.
    - This was the worst case for the DST
    - **However, we require no comparisons of the entire key**
  - So, again, a **benefit to RST is that the entire key does not need to be stored or compared**



- How can we improve tries?
  - ▶ If we can **reduce the heights** we will shorten the path length to the leaves
  - ▶ If we can **process multiple bits at a time** we can avoid bit operations
- We will examine a two variations that improve over the basic trie
  - ▶ **Multiway Trie** (in text as R-way Trie)
    - Allow more than 2 branches per node
  - ▶ **de la Brandeis Tree** (DLB)
    - Maintain fast lookups while reducing memory



- ▶ RST that we have seen considers the key 1 bit at a time.
  - This causes a **maximum height** in the tree of **b**, where b is the bits in the longest key stored.
  - A **search miss**, as explained in Proposition H on p. 743 of Sedgewick, will require on **average  $\lg_2 N$  nodes** to be examined (similar to BST).
- ▶ If we **considered m bits at a time**, then we could reduce the tree height.
  - **Maximum height** is now  **$b/m$**  since m bits are consumed at each level.
  - Let  $R = 2^m$ 
    - Average nodes on **search miss** is now  **$O(\log_R N)$** , since we branch in R directions at each node.





- ▶ Let's look at an example
  - Consider  $2^{20}$  (1 meg) keys of length 40 bits
    - Simple RST will have
      - > Worst Case height = 40
      - > Ave. nodes per miss =  $O(\log_2[2^{20}]) \approx 20$
    - Multiway Trie using 8 bits would have
      - > Worst Case height =  $40/8 = 5$
      - > Ave. nodes per miss =  $O(\log_{256}[2^{20}]) \approx 2.5$
- ▶ This is a considerable improvement
- ▶ Let's look at an example using character data (ex: a String)




### ► Consider a String: DATUM

- This is 40 bits, 8 bits per character
  - If we considered it 1 bit at a time, our tree would have a height of 40
  - However, if we can consider it 1 character (8 bits) at a time we can reduce the height to  $40/8 = 5$
- How can we do this?
  - Each bit has 2 possible values (1 or 0)
  - Each character has  $2^8 = 256$  possible values (ASCII)
  - Thus we need to somehow allow **256-way branching** at each node in our tree
  - How can we do this?
    - > Interesting question!



- ▶ Let's define a node in our tree to be an array of 256 pointers (references)
  - Each reference will point to a child node on the next level of the tree
  - Our array will be indexed on all of the possible character values in the ASCII set
    - Recall that ASCII values are simply integers from 0-255 which are interpreted as character values
    - We can use these values to index our array, giving us:



- ▶ Now the value of each character in the String will determine the index and branch to take
- ▶ Branching based on characters **reduces the height greatly**
  - If a string with  $K$  characters has  $n$  bits, it will have  $n/8$  ( $= K$ ) characters and therefore at most  $K$  levels
- ▶ Character / byte comps are simple, so, we **eliminate the bit comparison problem mentioned previously as item (2)**
- ▶ Thus, given  $N$  strings in a multiway trie, we have
  - **Theta(1) for a successful search or insert**
    - Theta( $K$ ) in terms of the key length
  - Theta( $\lg_R N$ ) node accesses for unsuccessful search



### ► So what is the catch (or cost)?

- **Memory**
  - Multiway Tries use considerably more memory than simple tries
- Each node in the multiway trie contains  **$R (= 2^m)$  pointers/references**
  - In example with ASCII characters,  $R = 256$
  - With simple words,  $R$  could be 26
- Many of these are unused, especially
  - During **common paths** (prefixes), where there is no branching (or "one-way" branching)
    - > Ex: through and throughout
  - At the **lower levels** of the tree, where previous branching has likely separated keys already
    - > Unique suffixes of words



- **Multiway Trie Structure**

- ▶ As mentioned, a node represents the possible characters at a given position in the key
  - This is represented by an **array of references**
  - Indices on array are the characters in the character set
  - A **non-null** value at an index indicates that a key exists with that char at that point in the key
  - Branching occurs when more than one of the pointers is non-null
- ▶ Each new key added will have some **prefix** represented by nodes already in the trie, followed by a **suffix** represented by new nodes



- ▶ Keys of different lengths are handled as follows:
  - In addition to the array of pointers, each node has a **field to indicate if a key terminates there**
    - It is set to **true** to indicate that the string up to that point is a key
    - It is set to **false** to indicate that the string up to that point is not a key
      - > Rather it is only a prefix to a key
- ▶ Note that keys are not directly stored in the trie at all
  - Rather each path from the root to a terminator node represents an individual key in the trie

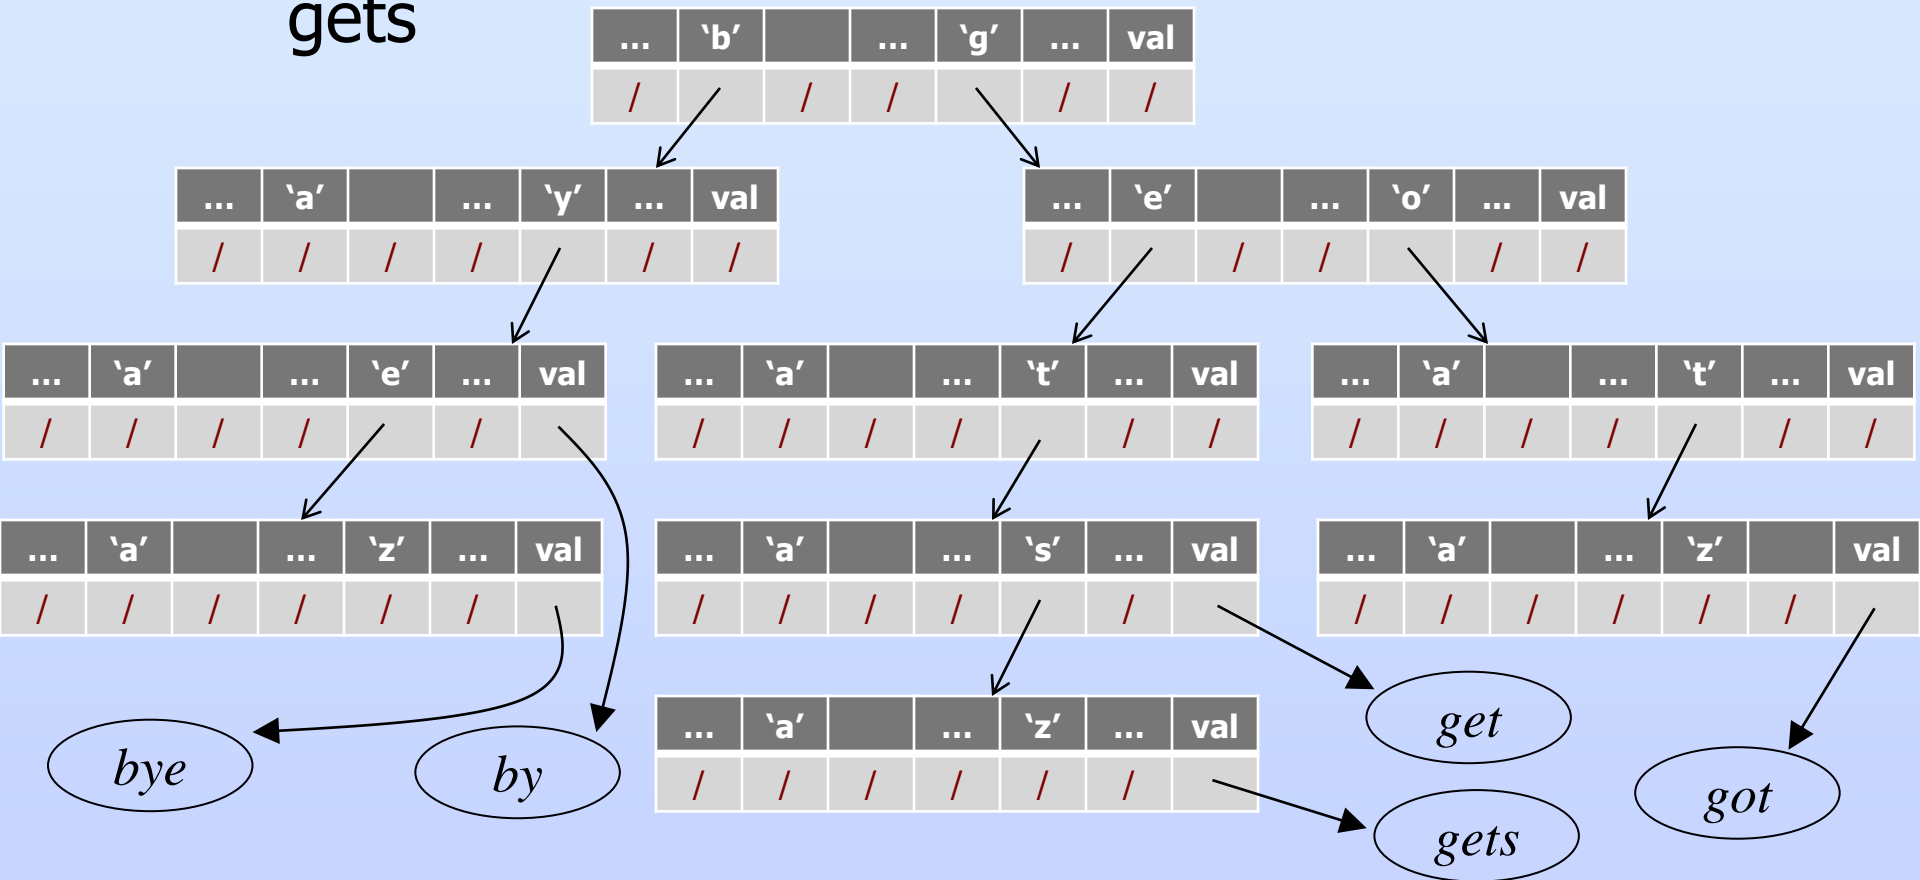


- ▶ Alternatively, if we make our trie **symbol table**, for each node we can have a reference to store the associated value
  - Recall that a symbol table maps **keys** to **values**
  - For a given node, if the value reference is non-null, then that node terminates a key, and the key (expressed as a path from the root to that node) is thus in the trie
    - Note that we need to proceed **one level down** from that last match
  - For a given node, if the value reference is null, then that node is within a key but does not terminate it, so the path from the root to that node is a **prefix of a key only**





- ▶ Ex: Consider strings: *bye*, *by*, *get*, *got* and *gets*



- ▶ See `TrieST.java` and `DictTest.java`



- Let's review again the 4 issues with DSTs:
  - 1) **Data is not sorted** – As with the simple RST a traversal here (done correctly) will produce the data in sorted order (look at figure in previous slide)
  - 2) **Bitwise operations are not easy** – comparing 8 bits (a char) at a time IS easy and efficient
  - 3) **What about keys of different lengths?** We just saw how to handle this with a termination field or value reference
  - 4) **Comparisons of the entire key** – these are not done at all



- Multiway tries solve these issues nicely
  - ▶ But with a **cost of a lot of memory**
  - ▶ This is especially evident when the symbol table is not very dense
    - i.e. most of the possible string keys do not actually exist in the symbol table
      - Yet we have references to all possible children in each node
    - Think about a dictionary of real words vs. the number of possible permutations of letters
  - ▶ Is there a way to keep (mostly) the "good" of a multiway trie while reducing memory?

