

Course Notes for
CS 1501
Algorithm Implementation

By
John C. Ramirez
Department of Computer Science
University of Pittsburgh



- These notes are intended for use by students in CS1501 at the University of Pittsburgh and no one else
- These notes are provided free of charge and may not be sold in any shape or form
- These notes are NOT a substitute for material covered during course lectures. If you miss a lecture, you should definitely obtain both these notes and notes written by a student who attended the lecture.
- Material from these notes is obtained from various sources, including, but not limited to, the following:
 - Algorithms in C++ by Robert Sedgewick
 - Algorithms, 4th Edition by Robert Sedgewick and Kevin Wayne
 - Introduction to Algorithms, by Cormen, Leiserson and Rivest
 - Various Java and C++ textbooks
 - Various online resources (see notes for specifics)



Improving Divide and Conquer Multiplication

- Previously we discussed two algorithms for multiplying N-bit integers:
 - Gradeschool, requiring $\Theta(N^2)$ time
 - Simple divide and conquer, also $\Theta(N^2)$
- Between the two, we would prefer gradeschool due to less overhead
 - However, maybe we can make the divide and conquer algorithm asymptotically better
 - Let's reconsider this algorithm and see how we can improve it
 - Let's start by considering the **recurrence relation** for it



Improving Divide and Conquer Multiplication

$$T(N) = 4T(N/2) + \Theta(N)$$

- ▶ We can try to reduce the amount of work in the current call
 - This could work, but will not in this case
 - This work sums to $\Theta(N^2)$ but so does the left part of the equation, so reducing it will keep the overall time at N^2
- ▶ We can try to make our subproblems smaller
 - Ex: $N/3$ or $N/4 \rightarrow$ but this would complicate our formula and likely require more subproblems
- ▶ We can try to reduce the number of subprobs
 - If possible, without changing the rest of the recur.



- Karatsuba's Algorithm

- If we can **reduce** the number of **recursive calls** needed for the divide and conquer algorithm, perhaps we can improve the run-time

- How can we do this?

- Let's look at the equation again

$$\begin{array}{ccccccc}
 XY & = & 2^N X_H Y_H & + & 2^{N/2} (X_H Y_L + X_L Y_H) & + & X_L Y_L \\
 & & (M_1) & & (M_2) & (M_3) & (M_4)
 \end{array}$$

- Note that we don't really NEED M_2 and M_3 individually
 - All we need is the **SUM OF THE TWO, $M_2 + M_3$**
- If we can somehow derive this sum using only one rather than two multiplications, we can improve our overall run-time



More Integer Multiplication

- Now consider the following product:

$$(X_H + X_L) * (Y_H + Y_L) = X_H Y_H + X_H Y_L + X_L Y_H + X_L Y_L$$

> Note: This is the same as the original product but without any shifting of the high bits

- Using our M values from the previous slide, this equals

$$M_1 + M_2 + M_3 + M_4$$

- The value we want is $M_2 + M_3$, so define M_{23}

$$M_{23} = (X_H + X_L) * (Y_H + Y_L)$$

- And our desired value is $M_{23} - M_1 - M_4 = M_2 + M_3$

- Ok, all I see here is wackiness! How does this help?

– Let's go back to the original equation, and plug back in

$$\begin{aligned} XY &= 2^N X_H Y_H + 2^{N/2} (X_H Y_L + X_L Y_H) + X_L Y_L \\ &= 2^N M_1 + 2^{N/2} (M_{23} - M_1 - M_4) + M_4 \end{aligned}$$

- Only 3 mults needed: M_1 , M_4 and M_{23}



More Integer Multiplication

- But will this cause other parts of the recurrence to increase?
 - Looking back, we see that M_{23} involves multiplying at most $(N/2)+1$ bit integers, so asymptotically it is the same size as our other recursive multiplications
 - We have to do some extra additions and two subtractions, but these are all $\Theta(N)$ operations
- Thus, we now have the following recurrence:
$$T(N) = 3T(N/2) + \Theta(N)$$
- This solves to $\Theta(N^{\lg 3}) \approx \mathbf{\Theta(N^{1.58})}$
 - Now we have an asymptotic improvement over the Gradeschool algorithm
 - > Still a lot of overhead, but for large enough N it will run faster than Gradeschool
- See
 - http://en.wikipedia.org/wiki/Karatsuba_algorithm
 - http://www.javamex.com/tutorials/math/BigDecimal_BigInteger_performance_multiply.shtml



► Practical Use?

- Hybrid algorithm that uses Gradeschool until large enough N (ex: ~ 3000 decimal digits or $\sim 3000 \times \lg_2 10$ bits) and then switches to Karatsuba

► Can we do even better?

- If we multiply the integers indirectly using the **Fast Fourier Transform** (FFT), we can achieve a run-time of **$\Theta(N[\lg N][\lg \lg N])$**
- This requires even larger numbers before it shows superiority (10s of thousands of decimal digits)
- Don't worry about the details of this algorithm
 - But if you are interested look at

http://en.wikipedia.org/wiki/Sch%C3%B6nhage-Strassen_algorithm



- How about integer powers: X^Y

- ▶ Natural approach: **simple for loop**

```
ZZ ans = 1; // assume ZZ is very large int
for (ZZ ctr = 1; ctr <= Y; ctr++)
    ans = ans * X;
```

- ▶ This seems ok – one for loop and a single multiplication inside – is it **linear**?

- ▶ Let's look more closely

- Total run-time is

- 1) Number of iterations of loop ***

- 2) Time per multiplication**



This assumes that the results as we multiply are still N-bit numbers – we will see how to ensure this soon.

Exponentiation

- We already know 2) since we just did it
 - Assuming GradeSchool, $\Theta(N^2)$ for N-bit ints
- How about 1)
 - It seems linear, since it is a simple loop
 - In fact, it is **LINEAR IN THE VALUE of Y**
 - However, our calculations are based on N, the **NUMBER OF BITS in Y**
 - What's the difference?
 - > We know an N-bit integer can have a value of up to $\approx 2^N$
 - > So **linear in the value of Y** is **exponential in the bits of Y**
 - Thus, the iterations of the for loop are actually $\Theta(2^N)$ and thus our total runtime is $\Theta(N^2 2^N)$
- This is RIDICULOUSLY BAD
 - > Consider $N = 512$ – we get $(512)^2(2^{512})$
 - > Just how big is this number?



Exponentiation

- Let's calculate in base 10, since we have a better intuition about size
 - Since every 10 powers of 2 is approximately 3 powers of ten, we can multiply the exponent by 3/10 to get the base 10 number
 - So $(512)^2(2^{512}) = (2^9)^2(2^{512}) = 2^{530} \approx 10^{159}$
 - Let's assume we have a 10 GHz machine (10^{10} cyc/sec)
 - This would mean we need 10^{149} seconds
 - $(10^{149}\text{sec})(1\text{hr}/3600\text{sec})(1\text{day}/24\text{hr})(1\text{yr}/365\text{days}) = (10^{149}/(31536000)) \text{ years} \approx 10^{149}/10^8 \approx 10^{141} \text{ years}$
 - This is ridiculous!!
- But we need exponentiation for RSA, so how can we do it more efficiently?



- ▶ How about a divide and conquer algorithm
 - Divide and conquer is usually worth a try
- ▶ Consider
$$X^Y = (X^{Y/2})^2 \text{ when } Y \text{ is even}$$
how about when Y is odd?
$$X^Y = X * (X^{Y/2})^2 \text{ when } Y \text{ is odd}$$
- ▶ Naturally we need a base case
$$X^Y = 1 \text{ when } Y = 0$$
- ▶ We can easily code this into a recursive function
- ▶ What is the run-time?



- ▶ Let's see...our problem is to calculate the exponential X^Y for X and Y
 - So we have a recursive call with an argument of $\frac{1}{2}$ the original size, plus a multiplication (again assume we will use GradeSchool)
 - **We'll put the multiplication time back in later**
 - **For now let's determine the number of function calls**
 - How many times can we divide Y by 2 until we get to a base case?
 - Since Y is a N -bit integer, it could be up to 2^N
 - Thus, we will start at 2^N



Exponentiation

$$\text{Step 0: } 2^N = Y$$

$$\text{Step 1: } 2^{N-1} = Y/2^1$$

$$\text{Step 2: } 2^{N-2} = Y/2^2$$

...

$$\text{Step N: } 2^{N-N} = Y/2^N = 1$$

- ▶ How many total steps?

- $N+1 = \lg_2(Y) + 1$

- ▶ Thus, the number of recursive calls is **logarithmic in Y** and **linear in N**

- Compare this to the for loop version

- ▶ Since we have one or two mults per call, we end up with a total runtime of $\Theta(N^2 * N) = \mathbf{\Theta(N^3)}$



- ▶ This is an AMAZING improvement
 - Consider again $N = 512$
 - $N^3 = 134217728$ – less than a billion
 - On a 10GHz machine this would take **less than a second**
 - Remember that or naïve algorithm required 10^{141} **years**
 - **Think about that difference!**
- ▶ **But is this result actually correct?**
 - Let's think about our result value



- ▶ Note that the power function can create enormous numbers
 - If X is N bits, X^2 is $2N$ bits, X^3 is $3N$ bits and so on
 - This increases the time required for the next multiplication and causes a lot of overhead for memory allocation
 - In the end our run-time will be dominated by multiplication itself, once the numbers get HUGE^{HUGE}
 - We also could not fit these numbers in memory
 - In practice (ex: for encryption) we perform the power operation **modulo some other value**
 - The final result must therefore be less than the modulo value, keeping the run-time for multiplication and the memory overhead in check
 - See Power.java



► Can we improve even more?

- Well removing the recursion can always help
- If we start at X , then square repeatedly, we get the same effect as the recursive calls
- Square at each step, and also multiply by X if there is a 1 in the binary representation of Y (from left to right)
- Ex: $X^{45} = X^{101101} =$

$1, X$	X^2	X^4, X^5	X^{10}, X^{11}	X^{22}	X^{44}, X^{45}
1	0	1	1	0	1
- Same idea as the recursive algorithm but building from the "bottom up"
- See Power.java

